

CSE 4705 - Artificial Intelligence

Multiple Regression

Joe Johnson M.S., Ph.D., A.S.A.

Department of Computer Science and Engineering
University of Connecticut



Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Training Set and Notation

Our dataset for the univariate linear regression case:

Training Set and Notation

Our dataset for the univariate linear regression case:

| Size in feet ² (x) | Price (\$) in 1000's (y) |
|-----------------------------------|------------------------------|
| 2104 | 400 |
| 1416 | 232 |
| 1534 | 315 |
| 852 | 178 |
| ... | ... |

$$f_{w,b}(x) = wx + b$$

Training Set and Notation

Now we have the following for our dataset:

Training Set and Notation

Now we have the following for our dataset:

Multiple features (variables)

| Size in feet ² | Number of bedrooms | Number of floors | Age of home in years | Price (\$) in \$1000's |
|---------------------------|--------------------|------------------|----------------------|------------------------|
| 2104 | 5 | 1 | 45 | 460 |
| 1416 | 3 | 2 | 40 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Training Set and Notation

Now, let's consider notation:

Training Set and Notation

Now, let's consider notation:

Multiple features (variables)

| | Size in feet ² | Number of bedrooms | Number of floors | Age of home in years | Price (\$) in \$1000's |
|-----|---------------------------|--------------------|------------------|----------------------|------------------------|
| | x_1 | x_2 | x_3 | x_4 | |
| | 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | 2 | 40 | 232 |
| | 1534 | 3 | 2 | 30 | 315 |
| | 852 | 2 | 1 | 36 | 178 |
| | ... | ... | ... | ... | ... |

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| 1416 | 3 | 2 | 40 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

$i=2$

Notation:

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | 2 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation:

- n = number of features

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | 2 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation:

- n = number of features
- m = number of samples

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | (2) | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation:

- n = number of features
- m = number of samples (or data objects)

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | 2 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation:

- n = number of features
- m = number of samples (or data objects) in our training set

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | 2 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation:

- n = number of features
- m = number of samples (or data objects) in our training set
- \vec{x}_j = vector of values for the j^{th} feature

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | 2 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation:

- n = number of features
- m = number of samples (or data objects) in our training set
- \vec{x}_j = vector of values for the j^{th} feature
- $\vec{x}^{(i)}$ = vector of features for the i^{th} training example

Training Set and Notation

Multiple features (variables)

| Size in feet ² | <u>Number of bedrooms</u> | <u>Number of floors</u> | <u>Age of home</u> in years | Price (\$) in \$1000's |
|---------------------------|---------------------------|-------------------------|-----------------------------|------------------------|
| x_1 | x_2 | x_3 | x_4 | |
| 2104 | 5 | 1 | 45 | 460 |
| i=2 | 1416 | 3 | 2 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

Notation:

- n = number of features
- m = number of samples (or data objects) in our training set
- \vec{x}_j = vector of values for the j^{th} feature
- $\vec{x}^{(i)}$ = vector of features for the i^{th} training example
- $x_j^{(i)}$ = value of feature j for the i^{th} training example

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Multiple Regression Model

Model Notation:

Multiple Regression Model

Model Notation:

- Model: $f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$

Multiple Regression Model

Model Notation:

- Model: $f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$
- Model parameters:

Multiple Regression Model

Model Notation:

- Model: $f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$
- Model parameters:
 - $\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$

Multiple Regression Model

Model Notation:

- Model: $f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$
- Model parameters:
 - $\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$ (row vector)

Multiple Regression Model

Model Notation:

- Model: $f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$
- Model parameters:
 - $\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$ (row vector)
 - b

Multiple Regression Model

Model Notation:

- Model: $f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$
- Model parameters:
 - $\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$ (row vector)
 - b
- Model (Vector Notation):

Multiple Regression Model

Model Notation:

- Model: $f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$
- Model parameters:
 - $\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$ (row vector)
 - b
- Model (Vector Notation): $f_{\vec{w}, b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Univariate Regression Example: House Prices Model

Recall gradient descent for the univariate case, $f_{w,b}(x)$:

Univariate Regression Example: House Prices Model

Recall gradient descent for the univariate case, $f_{w,b}(x)$:

- ① Initialize w, b to random values.

Univariate Regression Example: House Prices Model

Recall gradient descent for the univariate case, $f_{w,b}(x)$:

- ➊ Initialize w, b to random values.
- ➋ Repeat until convergence (simultaneous update):

Univariate Regression Example: House Prices Model

Recall gradient descent for the univariate case, $f_{w,b}(x)$:

- ➊ Initialize w, b to random values.
- ➋ Repeat until convergence (simultaneous update):

$$w := w - \alpha \frac{\partial J(w, b)}{\partial w}$$
$$b := b - \alpha \frac{\partial J(w, b)}{\partial b}$$

Univariate Regression Example: House Prices Model

Substituting the partials for the univariate case, we have:

Univariate Regression Example: House Prices Model

Substituting the partials for the univariate case, we have:

- ① Initialize w , b to random values.

Univariate Regression Example: House Prices Model

Substituting the partials for the univariate case, we have:

- ➊ Initialize w , b to random values.
- ➋ Repeat until convergence (simultaneous update):

Univariate Regression Example: House Prices Model

Substituting the partials for the univariate case, we have:

- ➊ Initialize w, b to random values.
- ➋ Repeat until convergence (simultaneous update):

$$w := w - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}}_{\frac{\partial J(w,b)}{\partial w}}$$

$$b := b - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})}_{\frac{\partial J(w,b)}{\partial b}}$$

Univariate Regression Example: House Prices Model

Now, for the multiple regression case, $f_{\vec{w}, b}(\vec{x})$, we have:

Univariate Regression Example: House Prices Model

Now, for the multiple regression case, $f_{\vec{w}, b}(\vec{x})$, we have:

- 1 Initialize $w_1, w_2, w_3, \dots, w_n, b$ to random values.

Univariate Regression Example: House Prices Model

Now, for the multiple regression case, $f_{\vec{w}, b}(\vec{x})$, we have:

- ➊ Initialize $w_1, w_2, w_3, \dots, w_n, b$ to random values.
- ➋ Repeat until convergence (simultaneous update):

Univariate Regression Example: House Prices Model

Now, for the multiple regression case, $f_{\vec{w}, b}(\vec{x})$, we have:

- ① Initialize $w_1, w_2, w_3, \dots, w_n, b$ to random values.
- ② Repeat until convergence (simultaneous update):

$$w_1 := w_1 - \alpha \frac{\partial J(\vec{w}, b)}{\partial w_1}$$

$$w_2 := w_2 - \alpha \frac{\partial J(\vec{w}, b)}{\partial w_2}$$

⋮

$$w_n := w_n - \alpha \frac{\partial J(\vec{w}, b)}{\partial w_n}$$

$$b := b - \alpha \frac{\partial J(\vec{w}, b)}{\partial b}$$

Univariate Regression Example: House Prices Model

Substituting the partials for the multiple regression case, we have:

Univariate Regression Example: House Prices Model

Substituting the partials for the multiple regression case, we have:

$$w_1 := w_1 - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_{\vec{w}, b}(\vec{x}^{(i)}) - y^{(i)}) \cdot x_1^{(i)}}_{\frac{\partial J(\vec{w}, b)}{\partial w_1}}$$

⋮

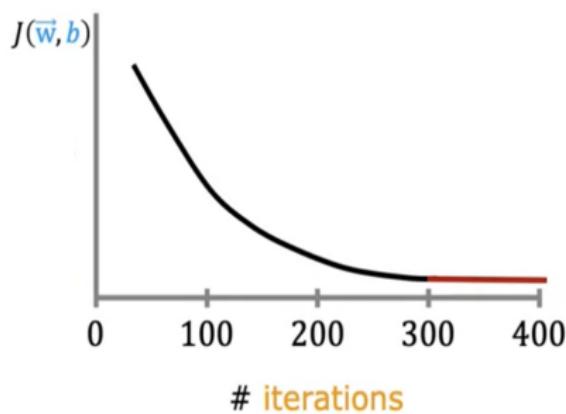
$$w_n := w_n - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_{\vec{w}, b}(\vec{x}^{(i)}) - y^{(i)}) \cdot x_n^{(i)}}_{\frac{\partial J(\vec{w}, b)}{\partial w_n}}$$

$$b := b - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (f_{\vec{w}, b}(x^{(i)}) - y^{(i)})}_{\frac{\partial J(\vec{w}, b)}{\partial b}}$$

Checking for Convergence

Make sure gradient descent is working correctly

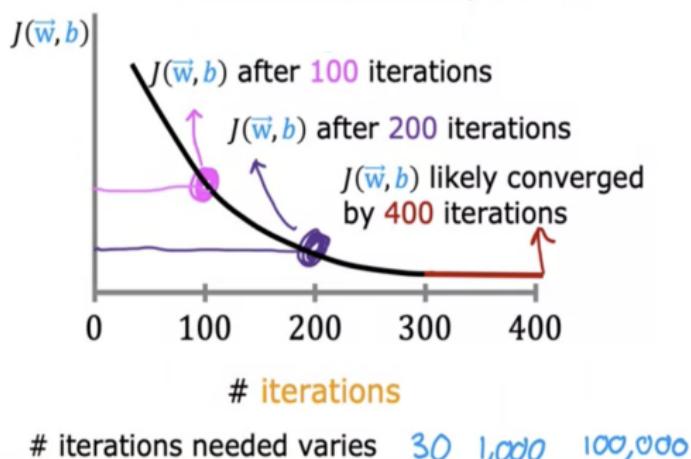
objective: $\min_{\vec{w}, b} J(\vec{w}, b)$



Checking for Convergence

Make sure gradient descent is working correctly

objective: $\min_{\vec{w}, b} J(\vec{w}, b)$ $J(\vec{w}, b)$ should decrease after every iteration



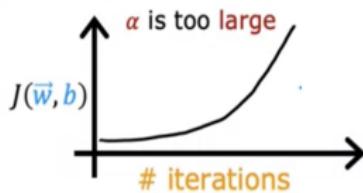
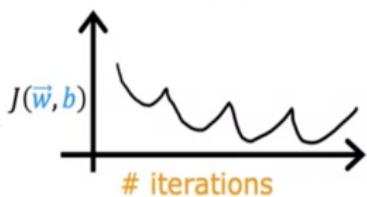
Automatic convergence test
Let ϵ "epsilon" be 10^{-3} .

If $J(\vec{w}, b)$ decreases by $\leq \epsilon$ in one iteration,
declare convergence.

(found parameters \vec{w}, b
to get close to
global minimum)

Checking Learning Rate

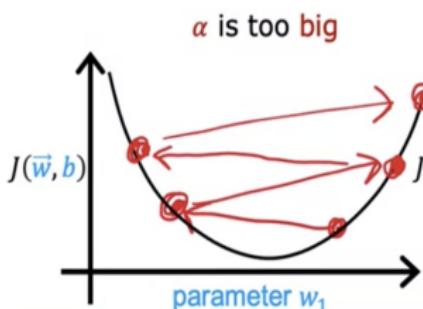
Identify problem with gradient descent



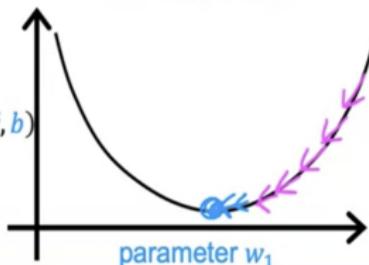
or learning rate is too large

$w_1 = w_1 + \alpha d_1$ ||
 use a minus sign
 $w_1 = w_1 - \alpha d_1$ ||

Adjust learning rate



Use smaller α



With a small enough α ,
 $J(\vec{w}, b)$ should decrease
 on every iteration

Learning Rate Optimization

Values of α to try:

... 0.001 0.003 $\xrightarrow{3X}$ 0.01 0.03 $\approx 3X$ 0.1 $\approx 3X$ 0.3 $\approx 3X$ 1 ...

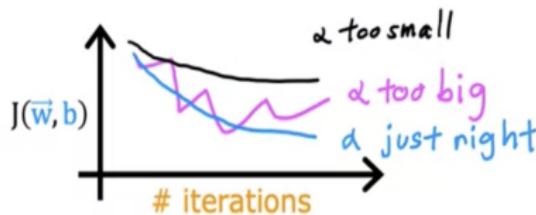
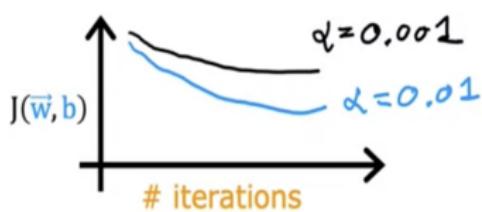


Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Feature Scaling

Feature and parameter values

$$\widehat{\text{price}} = w_1 x_1 + w_2 x_2 + b$$

x_1 : size (feet²) x_2 : # bedrooms
 range: 300 – 2,000 range: 0 – 5
 \downarrow \downarrow
 size #bedrooms large small

House: $x_1 = 2000$, $x_2 = 5$, $\text{price} = \$500k$ one training example

size of the parameters w_1, w_2 ?

$$w_1 = 50, \quad w_2 = 0.1, \quad b = 50$$

$$\widehat{\text{price}} = \underbrace{50 * 2000}_{100,000K} + \underbrace{0.1 * 5}_{0.5K} + \underbrace{50}_{50K}$$

$$\widehat{\text{price}} = \$100,050.5K = \$100,050,500$$

$$w_1 = 0.1, \quad w_2 = 50, \quad b = 50$$

small large

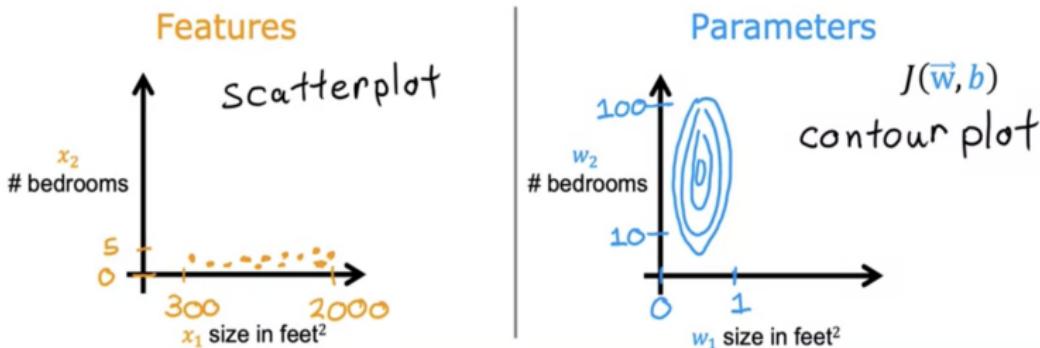
$$\widehat{\text{price}} = \underbrace{0.1 * 2000K}_{200K} + \underbrace{50 * 5}_{250K} + \underbrace{50}_{50K}$$

$$\widehat{\text{price}} = \$500k \text{ more reasonable}$$

Feature Scaling

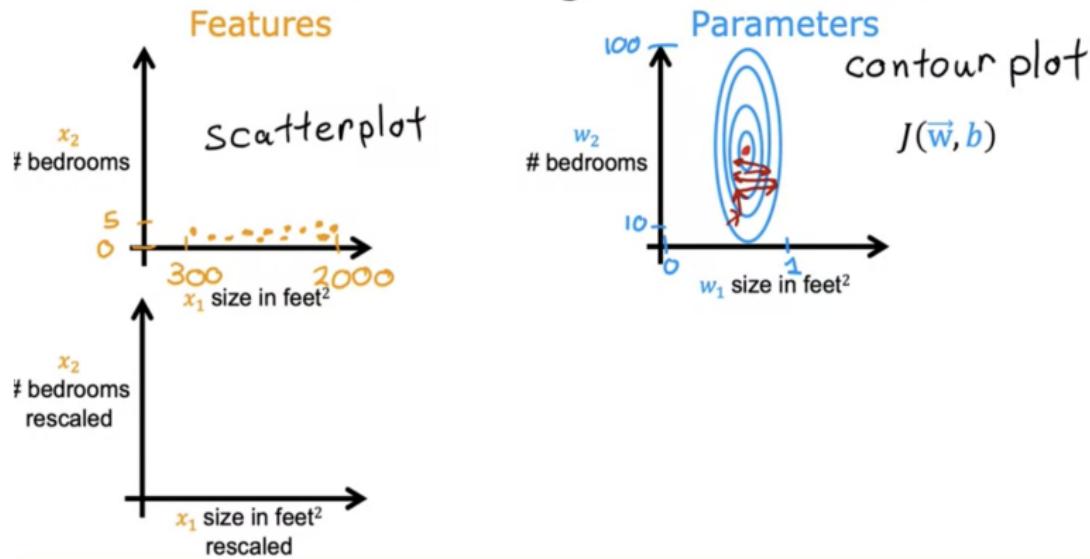
Feature size and parameter size

| | size of feature x_j | size of parameter w_j |
|---------------------------|-----------------------|-------------------------|
| size in feet ² | ↔ | ↔ |
| #bedrooms | ↔ | ↔↔ |



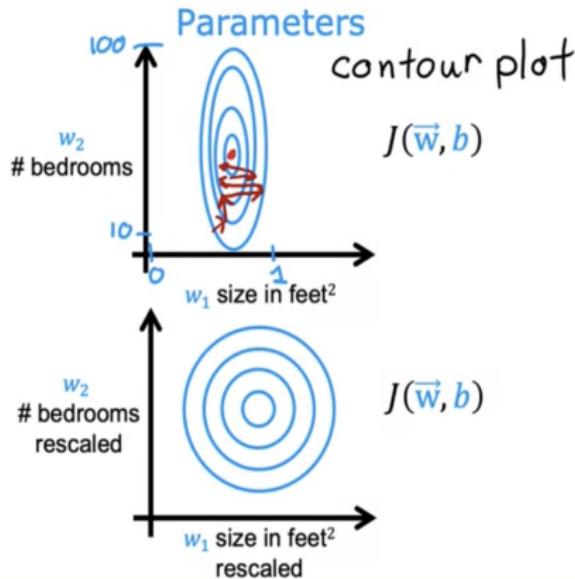
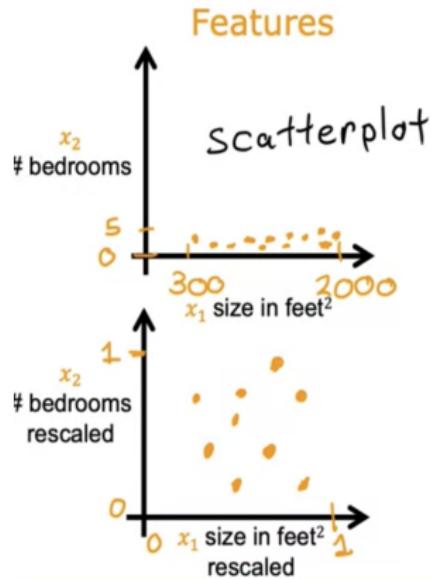
Feature Scaling

Feature size and gradient descent



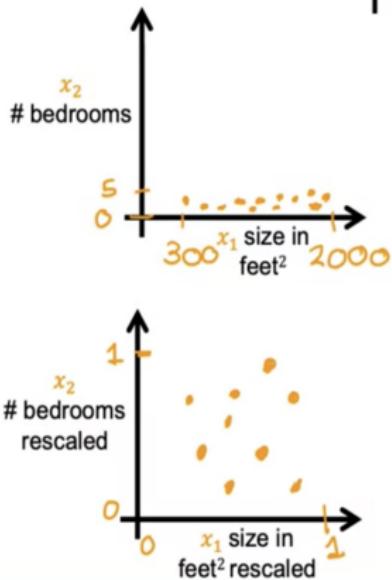
Feature Scaling

Feature size and gradient descent



Feature Scaling

Feature scaling



$$300 \leq x_1 \leq 2000 \quad 0 \leq x_2 \leq 5$$

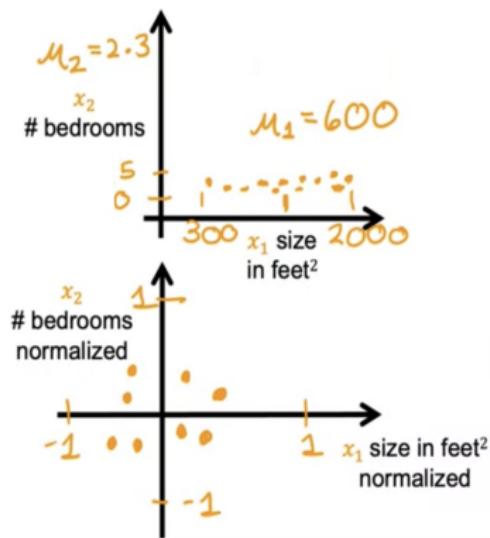
$$x_{1,scaled} = \frac{x_1}{\max} \quad x_{2,scaled} = \frac{x_2}{\max}$$



$$0.15 \leq x_{1,scaled} \leq 1 \quad 0 \leq x_{2,scaled} \leq 1$$

Feature Scaling

Mean normalization



$$300 \leq x_1 \leq 2000$$

$$x_1 = \frac{x_1 - \mu_1}{2000 - 300}$$

$$0 \leq x_2 \leq 5$$

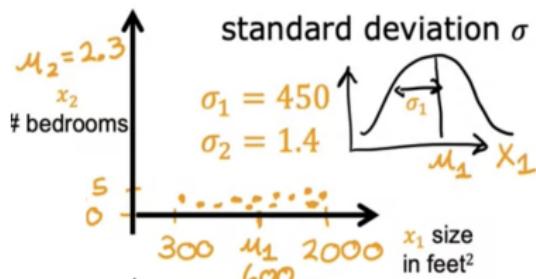
$$x_2 = \frac{x_2 - \mu_2}{5 - 0}$$

$$-0.18 \leq x_1 \leq 0.82$$

$$-0.46 \leq x_2 \leq 0.54$$

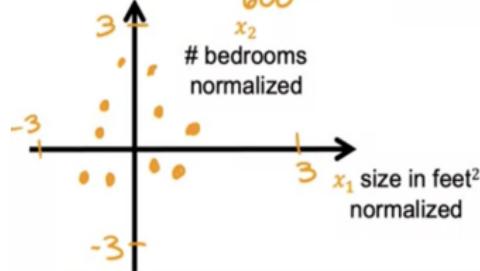
Feature Scaling

Z-score normalization



$$300 \leq x_1 \leq 2000 \quad 0 \leq x_2 \leq 5$$

$$x_1 = \frac{x_1 - \mu_1}{\sigma_1} \quad x_2 = \frac{x_2 - \mu_2}{\sigma_2}$$



$$-0.67 \leq x_1 \leq 3.1 \quad -1.6 \leq x_2 \leq 1.9$$

Feature Scaling

Feature scaling

aim for about $-1 \leq x_j \leq 1$ for each feature x_j

$$\left. \begin{array}{l} -3 \leq x_j \leq 3 \\ -0.3 \leq x_j \leq 0.3 \end{array} \right\} \text{acceptable ranges}$$

$0 \leq x_1 \leq 3$ Okay, no rescaling

$-2 \leq x_2 \leq 0.5$ Okay, no rescaling

$-100 \leq x_3 \leq 100$ too large → rescale

$-0.001 \leq x_4 \leq 0.001$ too small → rescale

$98.6 \leq x_5 \leq 105$ too large

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering**
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Feature Engineering - Example

Feature engineering

$$f_{\vec{w}, b}(\vec{x}) = \underline{w_1} \underline{x_1} + \underline{w_2} \underline{x_2} + b$$

frontage depth

$$\text{area} = \text{frontage} \times \text{depth}$$

$$x_3 = x_1 x_2$$

new feature

$$f_{\vec{w}, b}(\vec{x}) = \underline{w_1} \underline{x_1} + \underline{w_2} \underline{x_2} + \underline{w_3} \underline{x_3} + b$$



Feature Engineering - Definition

Feature Engineering:

Feature Engineering - Definition

Feature Engineering: is the use of intuition and problem domain knowledge

Feature Engineering - Definition

Feature Engineering: is the use of intuition and problem domain knowledge to design new features by transforming or combining original features in order to improve the predictive power of the model

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

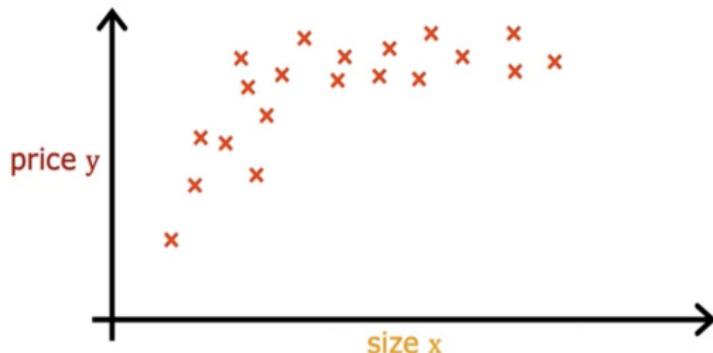
- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

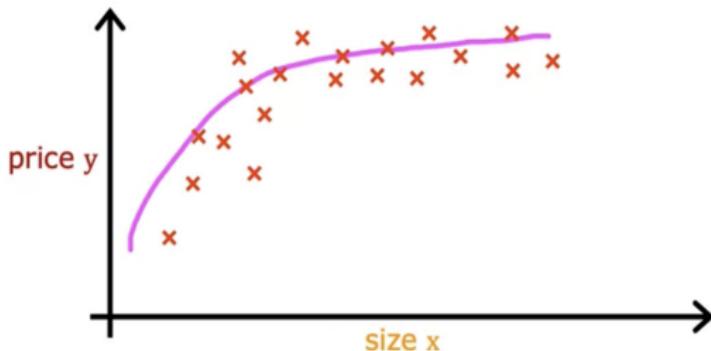
Polynomial Regression

Polynomial regression



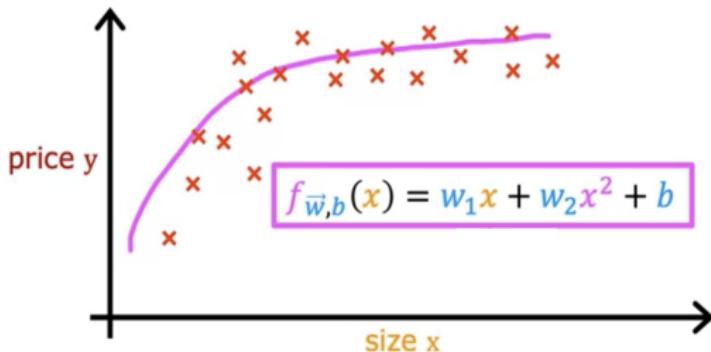
Polynomial Regression

Polynomial regression



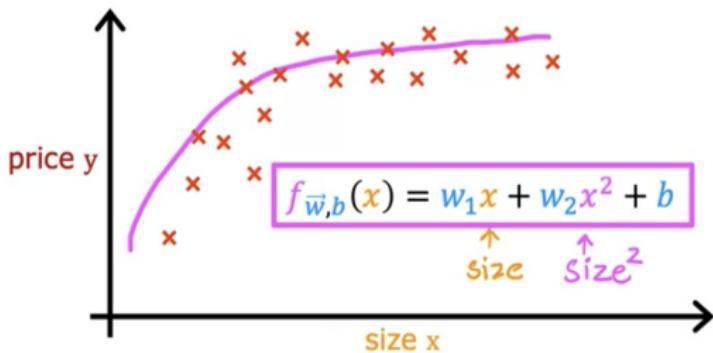
Polynomial Regression

Polynomial regression



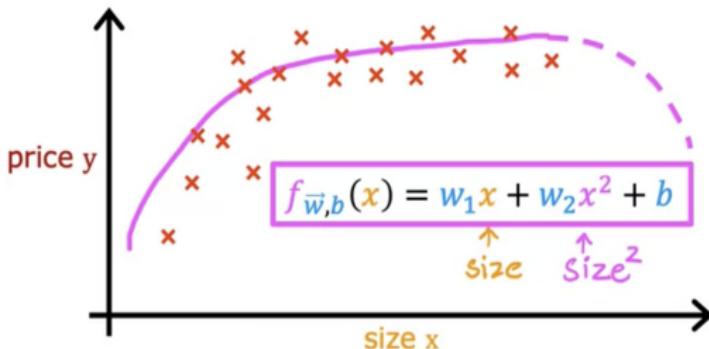
Polynomial Regression

Polynomial regression



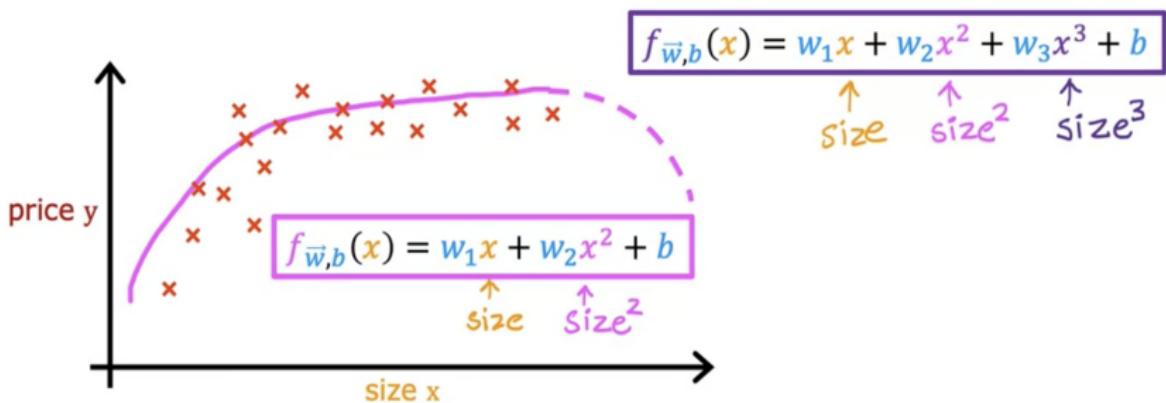
Polynomial Regression

Polynomial regression



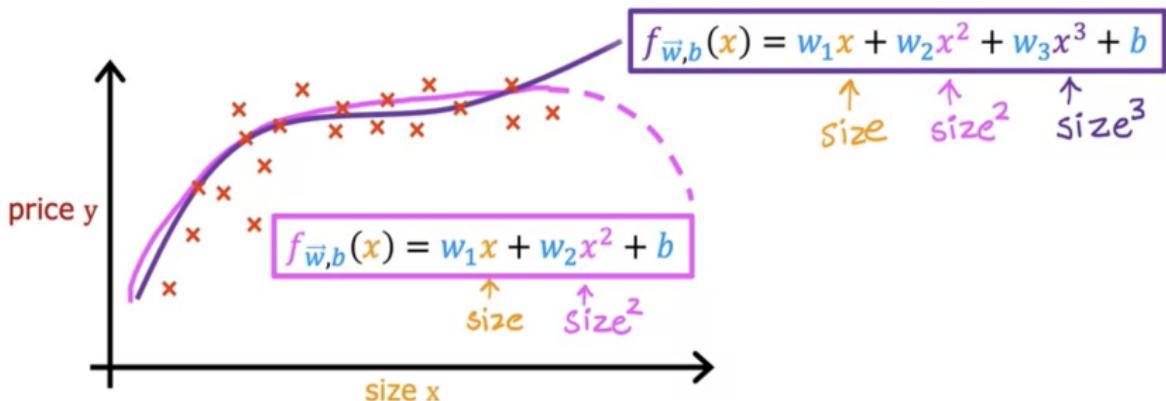
Polynomial Regression

Polynomial regression



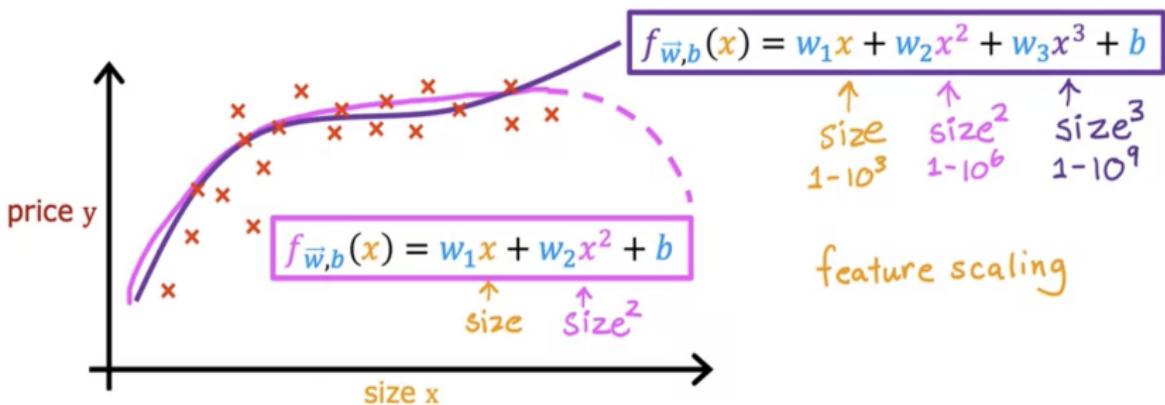
Polynomial Regression

Polynomial regression



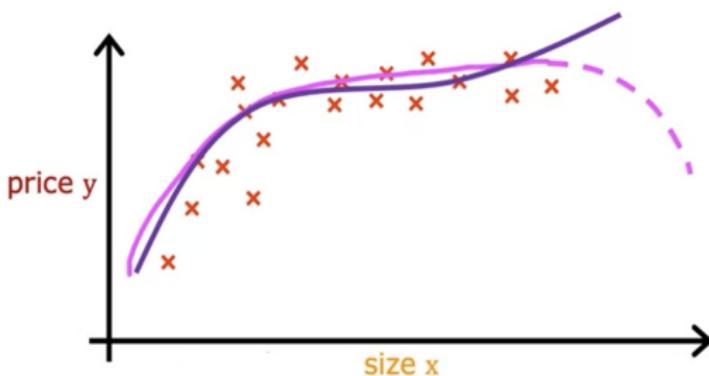
Polynomial Regression

Polynomial regression



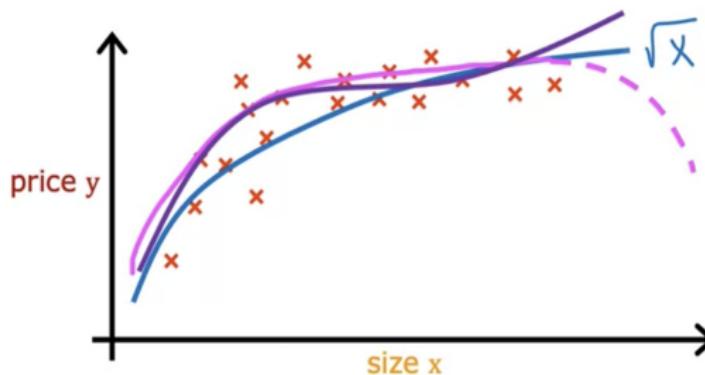
Polynomial Regression

Choice of features



Polynomial Regression

Choice of features



Polynomial Regression

Choice of features

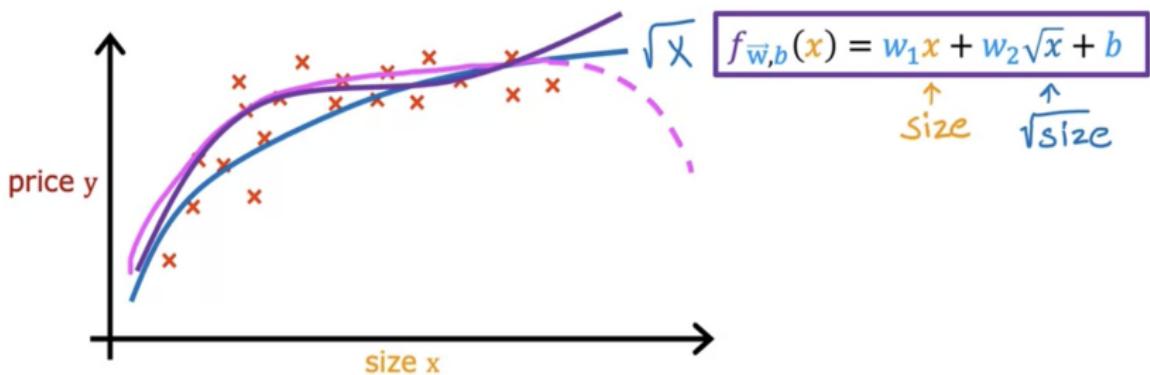


Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Table of Contents

1 Example - House Price Prediction

- Training Set and Notation

2 Model Development with Multiple Regression

- Multiple Regression Model
- Gradient Descent with Multiple Regression

3 Aspects of Model Selection

- Feature Scaling
- Feature Engineering
- Polynomial Regression

4 Addressing Computational Overhead

- Parallelization Using Vectorization

Vectorization

- Computation speed is an important issue in ML,

Vectorization

- Computation speed is an important issue in ML, especially as the size of our datasets gets large

Vectorization

- Computation speed is an important issue in ML, especially as the size of our datasets gets large both in terms of the number of samples and the number of features.

Vectorization

- Computation speed is an important issue in ML, especially as the size of our datasets gets large both in terms of the number of samples and the number of features.
 - We leverage the engineering of computing platforms to address this
-

Vectorization

- Computation speed is an important issue in ML, especially as the size of our datasets gets large both in terms of the number of samples and the number of features.
- We leverage the engineering of computing platforms to address this
 - **Numpy**.

Vectorization

Numpy:

Vectorization

Numpy:

```
w = np.array([1.0,2.5,-3.3])
b = 4
x = np.array([10,20,30])
```

Vectorization

Numpy:

```
w = np.array([1.0,2.5,-3.3])
b = 4
x = np.array([10,20,30])
```

- 0-index in code vs. 1-index in linear algebra

Vectorization

Consider once again our model:

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code?

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code? Here's one way:

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code? Here's one way:

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code? Here's one way:

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```

Unfortunately, this is clunky,

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code? Here's one way:

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```

Unfortunately, this is clunky, verbose,

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code? Here's one way:

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```

Unfortunately, this is clunky, verbose, brittle

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code? Here's one way:

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```

Unfortunately, this is clunky, verbose, brittle (we have to change code every time the number of features changes),

Vectorization

Consider once again our model:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

How do we implement the computation of $f_{\vec{w}, b}$ in code in code? Here's one way:

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```

Unfortunately, this is clunky, verbose, brittle (we have to change code every time the number of features changes), and slow.

Vectorization

Let's reconsider our model, this time using summation syntax:

Vectorization

Let's reconsider our model, this time using summation syntax:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

Vectorization

Let's reconsider our model, this time using summation syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \sum_{j=1}^n w_jx_i + b\end{aligned}$$

Vectorization

Let's reconsider our model, this time using summation syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \sum_{j=1}^n w_jx_i + b\end{aligned}$$

This may inspire us to use a for-loop:

Vectorization

Let's reconsider our model, this time using summation syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \sum_{j=1}^n w_jx_j + b\end{aligned}$$

This may inspire us to use a for-loop:

```
f = 0
for j in range(0,n):
    f = f + w[j] * x[j]
```

Vectorization

Let's reconsider our model, this time using summation syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \sum_{j=1}^n w_jx_j + b\end{aligned}$$

This may inspire us to use a for-loop:

```
f = 0
for j in range(0,n):
    f = f + w[j] * x[j]
```

This is an improvement, but it still computes the terms of f in series (and thus, is relatively slow).

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

$$f_{\vec{w}, b}(\vec{x}) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \vec{w} \cdot \vec{x} + b\end{aligned}$$

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \vec{w} \cdot \vec{x} + b\end{aligned}$$

This should inspire us to use vectorization syntax in our code:

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \vec{w} \cdot \vec{x} + b\end{aligned}$$

This should inspire us to use vectorization syntax in our code:

```
f = np.dot(w,x) + b
```

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \vec{w} \cdot \vec{x} + b\end{aligned}$$

This should inspire us to use vectorization syntax in our code:

```
f = np.dot(w,x) + b
```

One line of code!

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \vec{w} \cdot \vec{x} + b\end{aligned}$$

This should inspire us to use vectorization syntax in our code:

```
f = np.dot(w, x) + b
```

One line of code! More importantly, this is much faster.

Vectorization

Let's reconsider our model, yet again, but this time using vectorization syntax:

$$\begin{aligned}f_{\vec{w}, b}(\vec{x}) &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\&= \vec{w} \cdot \vec{x} + b\end{aligned}$$

This should inspire us to use vectorization syntax in our code:

```
f = np.dot(w, x) + b
```

One line of code! More importantly, this is much faster. Why?

Vectorization

Consider what happens when Python computes f using a for-loop:

Vectorization

Consider what happens when Python computes f using a for-loop:

```
for j in range(0,16):
    f = f + w[j] * x[j]
```

t_0
 $f + w[0] * x[0]$

t_1
 $f + w[1] * x[1]$

...

t_{15}
 $f + w[15] * x[15]$

Vectorization

Consider what happens when Python computes f using a for-loop:

```
for j in range(0,16):
    f = f + w[j] * x[j]
```

t_0
 $f + w[0] * x[0]$

t_1
 $f + w[1] * x[1]$

...

t_{15}
 $f + w[15] * x[15]$

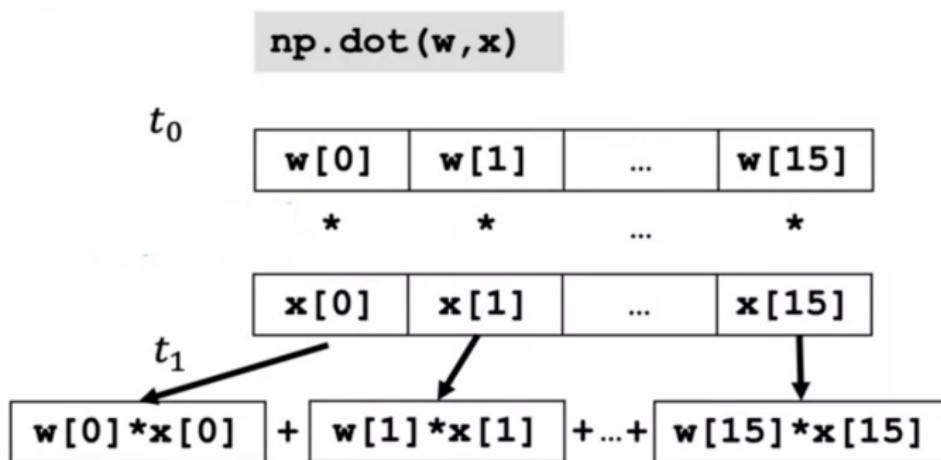
Each term is computed consecutively, (in series).

Vectorization

Now, consider what happens when Python computes f using vectorization:

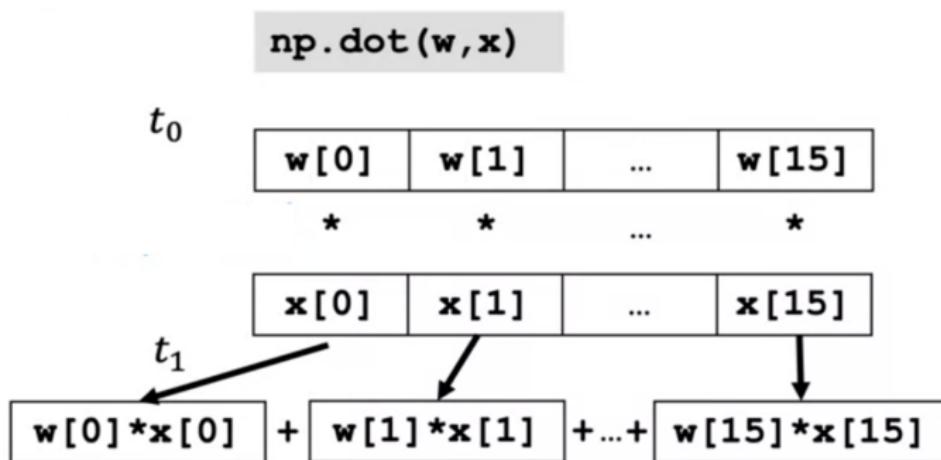
Vectorization

Now, consider what happens when Python computes f using vectorization:



Vectorization

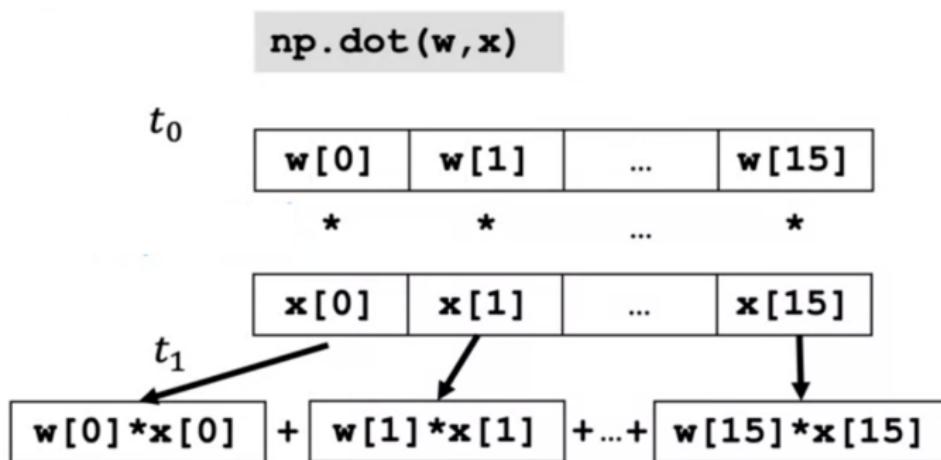
Now, consider what happens when Python computes f using vectorization:



Terms are computed in parallel,

Vectorization

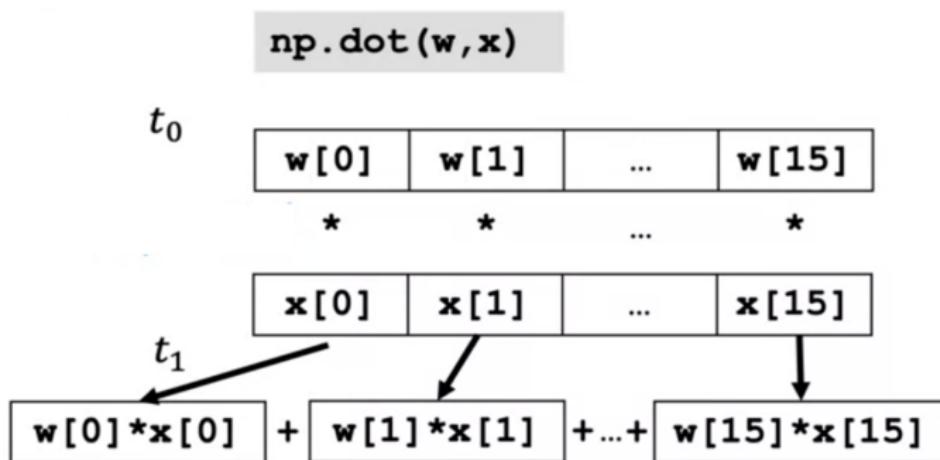
Now, consider what happens when Python computes f using vectorization:



Terms are computed in parallel, using parallel processing features of the underlying computer chip,

Vectorization

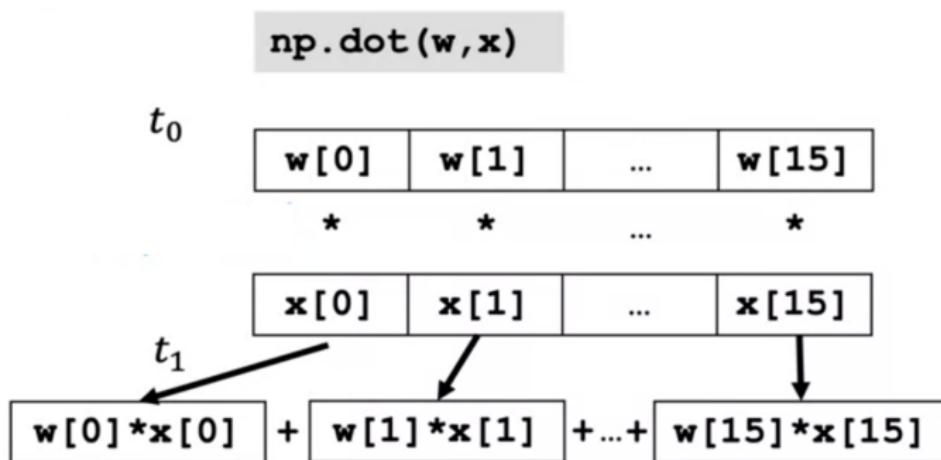
Now, consider what happens when Python computes f using vectorization:



Terms are computed in parallel, using parallel processing features of the underlying computer chip, leading to a dramatic improvement in performance.

Vectorization

Now, consider what happens when Python computes f using vectorization:



Terms are computed in parallel, using parallel processing features of the underlying computer chip, leading to a dramatic improvement in performance. (Scales well to large datasets.)

Vectorization

How do we apply parallelization in our gradient descent algorithm for multiple regression?

Vectorization

How do we apply parallelization in our gradient descent algorithm for multiple regression?

- Suppose we have the following:
 - $\vec{w} = (w_1, w_2, \dots, w_{16})$

Vectorization

How do we apply parallelization in our gradient descent algorithm for multiple regression?

- Suppose we have the following:
 - $\vec{w} = (w_1, w_2, \dots, w_{16})$ - vector of w parameters

Vectorization

How do we apply parallelization in our gradient descent algorithm for multiple regression?

- Suppose we have the following:
 - $\vec{w} = (w_1, w_2, \dots, w_{16})$ - vector of w parameters
 - $\vec{d} = (d_1, d_2, \dots, d_{16})$

Vectorization

How do we apply parallelization in our gradient descent algorithm for multiple regression?

- Suppose we have the following:

- $\vec{w} = (w_1, w_2, \dots, w_{16})$ - vector of w parameters
- $\vec{d} = (d_1, d_2, \dots, d_{16})$ - vector of partial derivatives wrt w 's

Vectorization

How do we apply parallelization in our gradient descent algorithm for multiple regression?

- Suppose we have the following:
 - $\vec{w} = (w_1, w_2, \dots, w_{16})$ - vector of w parameters
 - $\vec{d} = (d_1, d_2, \dots, d_{16})$ - vector of partial derivatives wrt w 's
- We store these vectors in numpy arrays:

Vectorization

How do we apply parallelization in our gradient descent algorithm for multiple regression?

- Suppose we have the following:
 - $\vec{w} = (w_1, w_2, \dots, w_{16})$ - vector of w parameters
 - $\vec{d} = (d_1, d_2, \dots, d_{16})$ - vector of partial derivatives wrt w 's
- We store these vectors in numpy arrays:

```
w = np.array([0.5, 1.3, ... 3.4])  
d = np.array([0.3, 0.2, ... 0.4])
```

Vectorization

- Non-vectorized approach using a for-loop (poor performance):

$$w_1 = w_1 - 0.1d_1$$

$$w_2 = w_2 - 0.1d_2$$

⋮

$$w_{16} = w_{16} - 0.1d_{16}$$

```
for j in range(0,16):  
    w[j] = w[j] - 0.1 * d[j]
```

Vectorization

- Vectorized approach:

$$\mathbf{w} = \mathbf{w} - 0.1 * \mathbf{d}$$

Vectorization

- Vectorized approach:

$$\mathbf{w} = \mathbf{w} - 0.1 * \mathbf{d}$$

- Much better performance

Vectorization

- Vectorized approach:

$$\mathbf{w} = \mathbf{w} - 0.1 * \mathbf{d}$$

- Much better performance
- Less code

Vectorization

- Vectorized approach:

$$\mathbf{w} = \mathbf{w} - 0.1 * \mathbf{d}$$

- Much better performance
- Less code
- Flexible - handles vectors of different sizes with no change to code